Student Project FAQ

1. What does this project do?

The Self Updating Documentation Assistant reads your TXT, Markdown and PDF documentation files, splits them into small passages, and answers questions by returning the passage that best matches your wording.

2. Which libraries are required?

Streamlit builds the web page. scikit-learn provides the TF-IDF vectorizer and cosine similarity. pypdf extracts text from PDF files. Nothing else is needed to run the app, and no API key or paid service is used.

3. Does the app need an internet connection?

No. After the packages are installed once, the whole application works offline on a normal laptop, because TF-IDF is pure local computation.

4. How do I start the application on Windows?

Open a terminal in the project folder, activate the virtual environment with .venv\Scripts\activate, then run the command streamlit run app.py. The browser opens at localhost port 8501 automatically.

5. How are answers generated?

The answer is extractive. The system selects the sentences from your own documents that have the highest cosine similarity with the question, so it quotes your documentation and never invents new facts.

6. What should I do when no answer is found?

Check that the document really covers the topic, use words that appear in the document, and lower the minimum similarity score in the sidebar.

7. How do I reset the search index?

Open the sidebar and press the Rebuild index button. This deletes the saved index files and processes every document again from scratch.